



Defense Logistics Agency (DLA) J6 Enterprise AI and Analytics Center of Excellence (AIACOE)

Request for Information (RFI): Cloud-Based LLM, RAG, and Agentic AI Support

June 23, 2026

Submitted By: C5MI Insight LLC
UID: L200000391

Primary Point of Contact: Garth Sanginiti, Chief Growth Officer
Address: 120 5th Street N, Jacksonville Beach, Florida, 32250
Email: garth.sanginiti@c5mi.com
Phone: +1 (636) 448-4125

Contractual Point of Contact: Megan Manley, Senior Contracts Manager
Address: 120 5th Street N, Jacksonville Beach, Florida, 32250
Email: megan.manley@c5mi.com
Phone: +1 (304) 268-1922

This proposal includes data that shall not be disclosed outside the Government and shall not be duplicated, used, or disclosed-in whole or in part-for any purpose other than to evaluate this proposal. If, however, a contract is awarded to this offeror as a result of-or in connection with-the submission of this data, the Government shall have the right to duplicate, use, or disclose the data to the extent provided in the resulting contract. This restriction does not limit the Government's right to use information contained in this data if it is obtained from another source without restriction. The data subject to this restriction are contained on the pages within this proposal.

Table of Contents

1.0 Cover Letter.....	ii
2.0 Executive Summary	1
3.0 Market Research Questions	1
3.1 Previous Enterprise Cloud / GenAI Experience	1
3.2 LLM and Foundation Model Experience	3
3.3 Agentic AI Experience.....	5
3.4 Data, Security, and Responsible AI	6
3.5 Integration and Enterprise Architecture.....	8
3.6 MLOps / LLMOps / Platform Operations	9
3.7 Implementation and Risk	10
3.8 Workforce, Pricing, and Delivery Model.....	11
3.9 Optional Government-Specific Questions	13
4.0 Recommendations.....	14
5.0 Conclusion	16

1.0 Cover Letter

June 23, 2026

TO:

Joseph Forgach, Contracting Officer (CO)

joseph.forgach@dla.mil

Re: DLA J6 AIACOE RFI for Cloud-Based Language Learning Model (LLM), Retrieval-Augmented Generation (RAG), AND Agentic Artificial Intelligence (AI) Support

Dear Mr. Forgach and other response reviewers and stakeholders,

Thank you for the opportunity to respond to the DLA RFI for Cloud-Based LLM, RAG, and Agentic AI Support. C5MI is confident that our demonstrated experience delivering secure, scalable enterprise AI capabilities in complex Federal environments, combined with our strategic partnership, positions our team to provide DLA with production-ready solutions that accelerate mission outcomes while maintaining strict compliance with federal security and Responsible AI standards.

Company Information:

Company Name: C5MI Insight LLC (C5MI)

Business Size and Socioeconomic Classification: Large Business (Other than Small); qualified as Small Business under JETS 2.0

Primary Point of Contact:

Garth Sanginiti, Chief Growth Officer

120 5th Street N, Jacksonville Beach, Florida 32250

Email: garth.sanginiti@c5mi.com

Phone: +1 (636) 448-4125

Teaming Structure

C5MI will serve as the Prime Contractor, leading an integrated Contractor Team Arrangement (CTA) designed to deliver secure, scalable, and production-ready AI capabilities aligned with DLA mission systems.



C5MI serves as an operational architect and system integrator. C5MI has **deployed SAP WMS solutions across 100+ DLA sites** and understands the security, compliance, and integration requirements for DLA enterprise systems. We are both **Google and SAP Partners** focused on "strengthening the digital core" by applying expertise in SAP modernization and Industry 4.0 to Google's cloud infrastructure, translating advanced technical capabilities into operational outcomes for defense and industrial supply chains. C5MI is DLA's SAP System Integrator and Managed Service Provider for **DLA - WMS, EBS and FLIS** Environments and specializes in operationalizing IT and AI. With SAP's recent pivot to helping its customers build and deploy "Autonomous Enterprise," C5MI is uniquely positioned in integration complex SAP Supply Chain and warehouse operations with IOT, Automation as well as integrating Google AI with DLA's business operations at scale!

KNEXUS **Knexus: 20 Years of AI Expertise & Proven DLA Track Record:** Knexus combines two decades of advanced AI expertise with deep Google Cloud Platform (GCP) engineering. We designed, secured, and deployed GAILFORCE, the first generative AI prototype to transition into DLA's GCP production environment (Fall 2025) and continue to sustain and enhance the application for agency-wide use. Knexus is also delivering DLA's first Data Fabric and Control Towers, unifying disconnected datasets (e.g., WMS, EBS, FLIS) into centralized hubs. This platform leverages GenAI and agentic workflows to enable automated alerting, real-time visualization, and natural language querying for mission-critical supply chain decisions.

A **Google Premier Partner** holding specialized expertise badges in **Gemini for Government, AI/ML, and Data Analytics**, Knexus was recognized as the 2025 and 2026 Google Cloud Business Applications Partner of the Year for Government.

 Google Cloud Google provides the foundational infrastructure for federal missions. Through [Gemini for Government](#) and [Vertex AI](#), they offer secure, scalable, and compliant agents capable of reasoning, planning, and executing tasks within strict regulatory boundaries. Google Public Sector is actively leading the effort to assist DLA in its AI Roadmap and AI deployment strategy. **Google Public Sector**, in close partnership with **KNEXUS** and **C5MI**, is also assisting DLA in building an enterprise-wide Data Fabric. The goal of the Data Fabric Project is to assist DLA in establishing a single source of truth across all disparate systems to serve as an enterprise backbone for all current and future AI initiatives.

Team C5MI understands the requirements and is fully capable of supporting the integration, modernization, and operational objectives across the systems, logistics environments, and task areas identified within the scope. Our team brings proven experience integrating complex enterprise platforms, data ecosystems, and mission-critical business processes to deliver secure, scalable, and mission-focused solutions.

This Contractor Team Arrangement (CTA) combines C5MI's proven DLA Warehouse Management System (WMS) and SAP platform delivery record with cutting-edge AI/ML capabilities, ensuring rapid deployment, operational resilience, and seamless integration with DLA's authoritative data sources and legacy business systems.

Relevant Cloud, LLM, RAG, and Agentic AI Capabilities

Our team delivers production-grade, secure enterprise AI platforms for Federal environments. We specialize in architecting RAG pipelines that eliminate model hallucinations by grounding LLM outputs to authoritative data sources with inline citations, advanced metadata filtering, and semantic document chunking. Our agentic AI solutions feature sandboxed tool execution, human-in-the-loop controls, and strict auditability, enabling autonomous multi-step workflows across Enterprise Resource Planning (ERP) ticketing systems, and data warehouses while preserving transaction integrity and approval chains. All deployments leverage model-agnostic architectures, FIPS 140-2 encryption, Customer-Managed Encryption Keys (CMEK), VPC isolation, and continuous MLOps/LLMOps pipelines to ensure compliance with DoD Responsible AI guidelines and operational continuity regardless of commercial model restrictions.

Our response provides comprehensive information addressing all RFI questions and demonstrates our commitment to delivering innovative, secure, and sustainable AI capabilities that accelerate DLA's mission. We look forward to supporting the DLA J6 AIACOE in establishing a transformational AI and analytics platform that enhances decision-making, operational efficiency, and mission readiness across the enterprise.

Respectfully submitted,

Donald Young
Chief Executive Officer
C5MI

2.0 Executive Summary

C5MI leads a team that includes Google and Knexus (herein after “Team C5MI”) to deliver cloud-based Large Language Model (LLM), Retrieval-Augmented Generation (RAG), and agentic AI capabilities for the Defense Logistics Agency. Our team combines direct DLA platform experience with enterprise AI implementation capabilities. C5MI has deployed SAP WMS solutions across 100+ DLA sites and understands the security, compliance, and integration requirements for DLA enterprise systems.



Enterprise AI Implementation. We design RAG pipelines that connect LLMs to authoritative Government data sources. Our implementations use semantic document chunking, metadata filtering, and citation generation to link every LLM response to specific source documents. When the knowledge base does not contain requested information, the system states that the data is unavailable rather than generating speculative answers. This approach ensures Government personnel receive verifiable outputs suitable for operational decisions.

Agentic AI with Operational Controls. Our agentic AI solutions execute multi-step workflows across ERP systems, ticketing platforms, and data warehouses. All agent actions run through API gateways with predefined schemas. We implement human-in-the-loop approval requirements for actions with operational or financial impact, rollback controls for failed transactions, and complete audit logging for all agent activities. This preserves existing approval chains and transaction integrity.

Model-Agnostic Architecture. We confirmed our understanding that Anthropic models are prohibited for this effort. Our architecture uses an abstraction layer that separates application logic from specific model implementations. This design allows DLA to switch between foundation models without modifying core workflows or data pipelines. The platform supports managed proprietary models, open-source models, and self-hosted deployments. If DoD restricts additional commercial models in the future, DLA operations continue without interruption.

Security and Compliance. All deployments use FIPS 140-2 validated encryption, Customer-Managed Encryption Keys for Government cryptographic control, and VPC Service Controls for workload isolation. We log every prompt, model call, tool execution, and administrative action to tamper-evident audit repositories. Our MLOps pipelines version-control prompts and system instructions alongside code, and automated evaluation gates validate changes before production deployment.

Implementation and Knowledge Transfer. We use Agile and DevSecOps methodologies with formal decision gates. Our timeline delivers pilot capability in months 1 to 2, production hardening in months 3 to 4, and scaled deployment in months 5 to 6. We provide Government teams with infrastructure-as-code, technical documentation, and hands-on training to enable independent platform operation after transition.

3.0 Market Research Questions

3.1 Previous Enterprise Cloud / GenAI Experience

1. **Describe your company’s experience designing, deploying, integrating, securing, and sustaining Enterprise cloud solutions that include large language models (LLMs), retrieval-augmented generation (RAG), agentic AI, machine learning, and analytics capabilities. Include the number of years of experience with each major capability proposed.**

Our combined C5MI Team offers over 20 years of experience designing, deploying, and sustaining enterprise-grade machine learning solutions, with 3 years of specific expertise in LLMs, RAG, and agentic AI. In partnership with Google Public Sector and C5MI, we integrate these capabilities with secure cloud architectures, including FedRAMP High and DoD Impact Level (IL) 4, 5, and 6 environments. Key implementations include GAILFORCE, a production-ready generative AI platform on Google Cloud that automates logistics data synthesis for the DLA, and Project EMMA, which modernizes DLA’s data infrastructure for secure RAG

querying. We have over 5 years of experience implementing RAG pipelines and deep semantic search, and 3 years of experience engineering stateful, goal-oriented agentic AI frameworks within zero-trust boundaries, ensuring all solutions strictly comply with DoD Responsible AI guidelines.

- 2. Describe your company's experience with Enterprise cloud AI services and related components, such as Vertex AI, Gemini models, Model Garden, Vector Search, RAG capabilities, agent frameworks, API management, container platforms, serverless services, data platforms, and workflow/orchestration tools. Include the number of years of implementation experience for each.**

Our team leverages three years of implementation experience with Google Cloud's enterprise AI suite, including Vertex AI, Gemini, and Model Garden. Team C5MI has integrated Vertex AI Vector Search and RAG capabilities for three years to support semantic search across diverse federal datasets. For API management and secure tool execution, we have utilized Google Cloud Apigee for over four years to enforce rigid OpenAPI schemas, complemented by zero-trust serverless architectures. Our containerized and serverless workloads (using Google Kubernetes Engine and Cloud Run) represent over six years of production experience. Furthermore, we leverage BigQuery (5 years) alongside Google Cloud Workflows for orchestrating complex AI-driven processes, ensuring a secure foundation that prevents data leakage.

- 3. Describe your company's experience delivering production AI/ML or generative AI capabilities for Federal agencies, especially DoD or other regulated environments. Include your three largest relevant clients and a short summary of the scope performed.**

Yes, Team C5MI successfully led the production implementation of GAILFORCE, a cloud-based generative AI solution for the DLA. Built on Google Cloud's secure infrastructure, the platform automates the synthesis of vast volumes of logistics data into actionable daily situation reports. This implementation achieved a significant reduction in data processing latency while completely eliminating manual data gathering bottlenecks. The scale of deployment spans multiple DLA directorates, securely processing terabytes of live mission-critical data for hundreds of active users. The period of performance commenced in 2024 and remains active, representing a mature, operational capability within the agency's environment.

- 4. Has your company led a Enterprise cloud-based AI/ML or generative AI implementation for a Federal agency that reached production or operational use? If yes, provide a summary of the project, mission outcomes achieved, scale of deployment, and period of performance.**

Yes. Team C5MI has successfully deployed GAILFORCE, a production-grade generative AI platform, for the Defense Logistics Agency (DLA). Built on Google Cloud's secure infrastructure, GAILFORCE automates the ingestion and synthesis of vast logistics datasets into actionable daily situation reports. This implementation has significantly reduced data processing latency and eliminated manual data-gathering bottlenecks. The deployment spans multiple DLA directorates, processing terabytes of mission-critical data for hundreds of active users, and has been fully accredited as a mature, operational DoD production system since 2024.

- 5. Describe your company's experience modernizing legacy applications or business processes by introducing Cloud-hosted AI capabilities, LLM-enabled user experiences, document understanding, semantic search, or AI-assisted workflow automation.**

Team C5MI modernizes legacy applications by wrapping existing transactional databases and document stores in secure, LLM-enabled user experiences. We leverage Google Cloud's document understanding and NLP capabilities to ingest, parse, and index legacy formats such as PDFs, technical manuals, and scanned records. Through semantic search, we enable operators to query complex institutional databases using natural language. For example, in our work with DLA, we integrated AI-assisted workflow automation to replace manual, multi-step administrative tasks with structured, automated checkpoints. This unlocks siloed data and enhances

decision-making without requiring a disruptive rewrite of the underlying legacy code, while maintaining strict human-in-the-loop controls.

6. Describe your company's experience supporting hybrid and multi-cloud environments where Enterprise Cloud AI services were integrated with existing enterprise systems, on-premises platforms, or other cloud providers.

Team C5MI supports resilient hybrid and multi-cloud environments by establishing unified data and AI orchestration layers. We integrate cloud-hosted AI services with existing on-premises platforms, legacy ERPs, and alternative cloud providers using secure, high-throughput API gateways. By utilizing Google Distributed Cloud and containerized deployment models, we execute analytical and ML workloads close to the data source. We implement secure data pipelines that aggregate multi-cloud datasets into unified vector stores, making them accessible to LLMs while maintaining strict security parity. This architecture ensures low-latency execution and complies with federal data residency requirements via centralized Identity and Access Management (ICAM).

7. Describe your company's experience supporting Government-owned data, Government-configured platforms, and Government-directed model/prompt/policy controls in enterprise AI environments.

Our model guarantees that the Government retains complete custody, sovereignty, and control over its data, platforms, and AI configurations. Team C5MI deploys all analytical workloads, model registries, and prompt libraries within tenant-isolated, Government-configured cloud enclaves (e.g., DoD IL 5/6). We enforce boundaries that prevent data leakage, ensuring Government-owned data is never used to train commercial models. Our administrative dashboards allow Government personnel to directly manage prompt libraries, system instructions, and safety parameters. All system interactions, model calls, and data access requests are fully logged to support federal auditing and compliance.

3.2 LLM and Foundation Model Experience

Please be advised that the use of Anthropic models (e.g., Claude) is strictly prohibited for this project. Please confirm your understanding of this restriction and describe the technical and procedural safeguards you have in place to ensure your proposed solution will remain compliant and will not utilize any Anthropic technologies.

Confirmed – Team C5MI and Google explicitly acknowledge and understand the restriction that the use of Anthropic models (e.g., Claude), services, APIs, or derivatives is strictly prohibited for this project. We have instituted rigid technical and procedural safeguards to ensure total compliance across the lifetime of the contract. To ensure continuous compliance, we employ technical safeguards that block unauthorized model endpoints, perform automated code and dependency scanning, and restrict model access through approved Google Cloud Vertex AI and authorized open-source frameworks. These controls are reinforced by strict configuration management, Government-approved change control processes, mandatory team training, and routine compliance audits to ensure only authorized AI technologies are used throughout contract performance.

8. Describe your company's experience selecting, evaluating, deploying, and tuning foundation models and LLMs for mission or enterprise use cases. Address your approach to choosing between managed proprietary models, open models, and model ensembles.

We utilize a rigorous, metrics-driven framework to select, evaluate, and tune models. We conduct trade-off analyses evaluating models against latency, context window efficiency, and licensing constraints. We deploy managed proprietary models (e.g., Google's Gemini) for high-capacity reasoning, while leveraging open models (e.g., Gemma) for local, containerized deployments. To optimize costs, we implement intelligent model routers that delegate routine tasks to smaller open models and escalate complex queries to larger ones. We apply advanced tuning methodologies including Parameter-Efficient Fine-Tuning (PEFT) and Low-Rank Adaptation

(LoRA) to align models with domain-specific terminology. Continuous re-evaluation against golden test datasets ensures safety and accuracy.

9. Describe your company's experience implementing LLM solutions for enterprise knowledge retrieval, conversational assistants, summarization, document analysis, code generation, classification, or decision support.

Google and Knexus possess deep operational experience building and sustaining enterprise-grade LLM. We design and deploy enterprise-grade LLM solutions tailored for knowledge retrieval, summarization, and decision support. We build conversational assistants that securely query multi-format institutional repositories using natural language. For document analysis, our automated pipelines extract key themes, action items, and metadata from massive, unstructured document sets. Our classification models automatically categorize transactional errors and supply chain disruptions, delivering context-aware recommendations. Additionally, we integrate specialized generative models within DevSecOps pipelines to automate legacy code translation and documentation. These capabilities are anchored to authoritative data sources to ensure all outputs are accurate, relevant, and fully auditable.

10. Describe your company's experience grounding LLMs on authoritative enterprise data to improve accuracy, relevance, and traceability of outputs. Include your experience with RAG architectures, vector stores, search, metadata filtering, document chunking, and citation generation.

Team C5MI eliminates model hallucinations by engineering high-fidelity RAG architectures that anchor outputs to verified enterprise data. Our ingestion pipelines segment complex documents by semantic boundaries, tagging each chunk with rich metadata. We utilize state-of-the-art vector databases and hybrid search engines that combine semantic vector embeddings with keyword search. This ensures highly relevant context retrieval across technical logistics terminology. To enforce absolute auditability, our RAG pipelines force the LLM to provide inline citations and direct links back to the exact source paragraph. If the authoritative data source does not contain the answer, the system is programmatically bound to declare the information unavailable rather than speculating.

11. Describe your company's methodology for evaluating model quality, including accuracy, hallucination rates, grounding fidelity, latency, cost, safety, task completion, and user adoption metrics.

Team C5MI's evaluation methodology assesses model quality across multiple dimensions. We deploy automated evaluation pipelines, such as Vertex AI Auto-SxS, to continuously score outputs against established golden datasets. Grounding fidelity is measured by calculating the percentage of facts directly traceable to source metadata, systematically minimizing hallucination. We track performance metrics like Time-to-First-Token (TTFT) and token consumption to optimize compute efficiency and lower total cost of ownership. Safety classifiers screen all outputs against federal guidelines. Finally, we monitor user adoption metrics, including copy-to-clipboard rates and explicit feedback loops, to drive continuous system refinement.

12. Explain your company's experience performing model benchmarking, A/B testing, red teaming, and continuous re-evaluation of LLM or agent performance in production.

Team C5MI executes structured model benchmarking, A/B testing, and adversarial red teaming to secure and optimize AI performance in production. Before deployment, we benchmark model variations against domain-specific baselines to measure reasoning and syntactic accuracy. In production, we establish canary deployments and shadow-routing pipelines to pass real-world data simultaneously through existing models and candidate variations. Our red-teaming protocols involve simulating advanced prompt injections and jailbreak attempts to verify safety guardrails. Continuous monitoring tools alert our DevSecOps team the moment live performance deviates from established statistical quality baselines.

13. Describe your company's approach to managing prompts, system instructions, reusable prompt assets, prompt versioning, and controlled prompt changes in production environments.

We treat prompts, system instructions, and few-shot context templates as core application code, enforcing strict engineering governance. Team C5MI implements centralized prompt registries that house all authorized system instructions in a secure repository. Every prompt modification is tracked using semantic versioning (SemVer) and linked to metadata specifying the validated models and safety configurations. Prompt updates must flow through a standardized DevSecOps CI/CD pipeline, requiring automated regression screening. Final promotion to production requires formal authorization from a Change Control Board, ensuring operational stability.

3.3 Agentic AI Experience

14. Describe your company's experience designing, deploying, and sustaining agentic AI solutions that can reason, plan, call tools, retrieve data, interact with APIs, and complete multi-step workflows.

Team C5MI team designs agentic AI solutions that transition systems from passive assistants into active, goal-oriented workflow execution engines. We build agents utilizing advanced reasoning techniques, such as Chain-of-Thought and ReAct frameworks, to decompose complex prompts into execution-ready sub-tasks. Agents interact with external systems via secure, managed APIs that enforce strict OpenAPI and JSON schemas. We design long-horizon execution capabilities that allow agents to query databases, apply conditional logic, and drive workflows to completion within strictly bounded permissions. Our sustainment model includes automated monitoring to ensure high reliability across all deployed agent pipelines.

15. Describe your company's experience implementing single-agent and multi-agent architectures for enterprise or Federal use cases. Include examples of mission outcomes achieved and lessons learned.

Team C5MI deploys both single-agent and multi-agent architectures tailored to complex, highly regulated operations. Single-agent deployments handle focused tasks like automated compliance analysis. For multi-layered missions, we implement collaborative multi-agent societies where specialized agents perform distinct roles (e.g., Retrieval, Validation, Synthesis). Lessons learned highlight the necessity of implementing mandatory maximum token boundaries and execution step caps to prevent runaway computational costs and optimize response latency.

16. Describe your company's experience orchestrating agents across enterprise systems such as ERPs, CRMs, ticketing platforms, document repositories, workflow tools, and data warehouses.

Team C5MI specializes in orchestrating AI agents across heterogeneous enterprise systems, including SAP, ServiceNow, and BigQuery. We design agent frameworks that interface seamlessly with ticketing software to interpret error logs, query historical workflows, and draft remediations. We enforce secure, cross-platform identity management utilizing role-based access controls (RBAC) and tightly constrained API gateways (e.g., Google Cloud Apigee). This ensures that every action taken across an ERP or document store is thoroughly authenticated, scoped, and audit-logged.

17. Describe your company's approach to determining when a use case should be implemented as a simple LLM workflow, a RAG solution, a rules-based automation, or a fully agentic solution.

We utilize a structured Use Case Classification Matrix to align requirements with the most cost-effective architecture. Team C5MI implements rules-based automation for high-volume, repetitive tasks. Simple LLM workflows are selected for static text manipulation. We deploy RAG when context-dependent knowledge retrieval is necessary. Fully agentic solutions are reserved for open-ended reasoning, multi-step planning, and workflows requiring real-time tool calling. This disciplined classification minimizes computational latency and optimizes platform costs.

18. Explain your company's approach to agent memory, session management, tool selection, human-in-the-loop approvals, rollback controls, and auditability of agent actions.

Team C5MI architecture implements a stateful framework managing agent memory via dual short-term session states and long-term external databases. For tool selection, we enforce strict semantic routing mechanisms backed by rigid OpenAPI schemas. If a tool call fails validation, the system triggers automated retry configurations or routes to exception handlers. We integrate mandatory state-saved approval checkpoints for high-consequence actions, ensuring no transaction is pushed to downstream systems without human validation.

19. Describe your company's experience evaluating AI agents, including measurement of task completion, tool-use quality, trajectory quality, policy compliance, and operational reliability.

Team C5MI utilizes a multi-tiered framework to measure task completion, tool-use quality, and trajectory efficiency. We analyze the agent's step-by-step reasoning trajectory to ensure efficient tool selection and avoidance of redundant actions. Tool-use quality is evaluated by monitoring syntax error rates and parameter accuracy. We rigorously test agent resilience against golden test suites representing real-world edge cases to measure operational reliability.

20. Describe your company's experience implementing agent guardrails to prevent unauthorized actions, invalid tool calls, policy violations, prompt injection, data leakage, or unsafe outputs.

We employ a defense-in-depth model that wraps agent execution in synchronous input and output guardrails. All incoming requests pass through input classification models (e.g., Model Armor) to filter malicious injections. Outbound responses are scanned in real-time to block PII leakage. Agents are sandboxed, possessing no raw network access, and interact exclusively via constrained API gateways. This prevents invalid tool calls and guarantees that an agent can never execute a command beyond its strictly assigned privileges.

3.4 Data, Security, and Responsible AI

21. Describe your company's approach to data governance for cloud-hosted LLM and agentic AI solutions. Include handling of structured and unstructured data, metadata, lineage, retention, and authoritative source management.

Our data governance framework manages structured operational logs and unstructured documents using enterprise platforms like BigQuery and Knowledge Catalog (formerly Dataplex). We isolate authoritative knowledge bases, ensuring models ground responses only on verified, signed data versions. Every data asset within our RAG pipelines is mapped with metadata identifying origin, classification level, and ingest timestamps. We maintain end-to-end lineage tracking for auditing, and our architectures enforce automated retention policies aligned with Federal and DoD records mandates.

22. Describe your company's experience implementing secure AI solutions in regulated Federal environments. Include identity and access management, encryption, key management, network isolation, logging, monitoring, and least privilege access controls.

We implement secure AI solutions adhering to DoD Impact Level 4, 5, and 6 standards. Team C5MI utilizes Virtual Private Cloud (VPC) Service Controls and software-defined perimeters to isolate execution workloads. We enforce strict ICAM utilizing RBAC and least-privilege routing. All data is protected using Customer-Managed Encryption Keys (CMEK). We route all system events into centralized, immutable logging engines, supported by real-time security monitoring to maintain ATO sustainment.

23. Describe your company's experience implementing responsible AI and safety controls for generative AI solutions. Include content moderation, harmful-content filtering, prompt and response inspection, bias testing, and escalation procedures.

Team C5MI builds and deploys systems in alignment with DoD Responsible AI (RAI) guidelines. We implement content moderation and real-time filtering using high-fidelity safety classifiers integrated into deployment pipelines. Every interaction undergoes synchronous inspection to redact harmful content or policy violations. During pre-deployment, model variations undergo bias and fairness testing. If an interaction triggers a safety violation, the system terminates the session and logs the incident, enabling rapid escalation to system administrators.

24. Explain your company's approach to protecting sensitive Government data from prompt injections, data exfiltration, unauthorized retention, model misuse, cross-tenant exposure, and insecure tool execution.

We implement a zero-trust model to guarantee data sovereignty. To protect against prompt injections, Team C5MI deploys synchronous input classification filters. Logical and physical tenant isolation ensures Government data is never cached, visible to other tenants, or used for commercial model training. We wrap autonomous operations in sandboxed enclaves that block raw network requests; agents interact exclusively via secure API management utilizing rigid OpenAPI schemas, ensuring tool execution remains bounded and secure.

25. Describe your company's experience implementing auditable controls for prompts, model calls, grounding sources, agent actions, and human approvals.

Team C5MI architecture establishes an immutable telemetry chain. We route all system events including prompts, instruction versions, grounding sources, and hyperparameters into a centralized logging engine. Every tool invocation is logged with transactional IDs to track the decision path. For high-consequence actions, we integrate mandatory approval gateways that pause execution until authorized. The system logs the official's identity, timestamp, and action, providing a seamless historical record for compliance.

26. Describe your company's experience supporting security and compliance requirements relevant to Federal and DoD environments, including approaches for boundary controls, workload isolation, data residency, and compliance evidence generation.

We support DoD Impact Level 4, 5, and 6 authorizations. Team C5MI utilizes strict VPC Service Controls to eliminate public internet entry vectors and shield environments. All data, models, and logs are bound exclusively to accredited Federal cloud data centers within CONUS boundaries, fulfilling data residency regulations. We leverage automated compliance tools that dynamically generate real-time security postures and configuration tracking to streamline continuous monitoring and ATO sustainment.

27. Describe your company's approach for ensuring that LLM or agent outputs remain explainable, attributable to source data where appropriate, and suitable for review by Government personnel before operational use.

Team C5MI forbids unverified speculation, anchoring responses to authoritative foundations. RAG pipelines force the LLM to provide explicit, inline citations and metadata links to source paragraphs. For agentic applications, the UI exposes the reasoning trajectory for operational review. Furthermore, outputs linked to downstream transactional processes are staged in a review interface, granting Government operators a comprehensive summary to confirm accuracy before final execution.

3.5 Integration and Enterprise Architecture

28. Describe your company's experience integrating Cloud-based AI capabilities with enterprise platforms such as SAP, ServiceNow, custom line-of-business applications, data lakes, warehouses, APIs, and legacy systems.

We modernize complex ecosystems by securely integrating enterprise cloud AI services into core transactional platforms. We leverage expertise in SAP ERP (S/4HANA and ECC) to pipeline data directly into cloud AI models. Team C5MI designs unified API connectors linking foundation models with platforms like ServiceNow and custom applications, enabling automated document parsing and error classification. Finally, we integrate deeply with modern data warehouses, allowing models to query multi-petabyte datasets securely.

29. Describe your company's approach to exposing AI services securely through APIs, user interfaces, chat experiences, workflow tools, or embedded enterprise applications.

Team C5MI expose AI services exclusively through a protected API management tier (e.g., Apigee). We implement tokenization, rate-limiting, and packet inspection. Whether accessed via chat, web UI, or embedded applications, the interface inherits the user's existing federal ICAM and RBAC controls. The application layer dynamically constrains data display and tool execution; guaranteeing operators interact only with features matching their explicit authorization.

30. Describe your company's experience with event-driven architectures, workflow orchestration, and tool invocation patterns used to connect AI agents with enterprise transactions and business processes.

We design decoupled, event-driven backbones linking AI agents to real-time transactions. Agents are triggered by enterprise events (e.g., logistics errors) via high-throughput message brokers, bypassing source system bottlenecks. For long-horizon execution, Team C5MI builds stateful orchestration loops enforcing logical branches and parallel coordination. Agents interact with line-of-business applications through strict JSON/RESTful tool-calling schemas managed via secure API layers.

31. Describe your company's experience implementing semantic search, enterprise search, document retrieval, and knowledge-grounding solutions across large and diverse repositories.

Team C5MI transforms fragmented, multi-format enterprise data into a unified, secure knowledge base. We deploy hybrid search engines combining semantic vector embeddings with keyword (BM25) search, allowing natural language queries against acronym-dense federal files. We design ingestion pipelines that parse documents based on semantic layout rather than arbitrary splits, storing them in vector databases with extensive metadata tagging. RAG pipelines act as strict boundaries, forcing the LLM to provide traceable links back to authoritative paragraphs.

32. Explain your company's approach to interoperability, portability, and avoidance of unnecessary vendor lock-in for AI, agent, and data architectures.

We decouple agent reasoning, prompt registries, and RAG pipelines from provider-specific APIs. By routing calls through a unified abstraction client, switching LLMs requires configuration adjustments rather than code rewrites. Team C5MI utilizes open-source containerization (Docker/Kubernetes) to ensure workloads migrate seamlessly across cloud or on-premises infrastructures. Data schemas are built on open formats, ensuring the Government retains full custody over assets and ensuring long-term operational continuity.

3.6 MLOps / LLMOps / Platform Operations

33. Describe your company's approach to MLOps and LLMOps for production Enterprise Cloud environments. Include CI/CD, model and prompt versioning, evaluation pipelines, testing, deployment promotion, rollback, monitoring, and cost management.

Team C5MI delivers a DevSecOps pipeline purpose-built for AI, managing prompts and configurations as code under semantic versioning. AI assets must pass automated evaluation pipelines against golden datasets before promotion. We utilize canary and shadow deployments to compare performance before broad rollout. If monitoring flags exceptions, automated rollback protocols revert to the system. We enforce rate-limiting, compute-quota bounding, and intelligent caching to maximize token efficiency and cost management.

34. Describe your company's experience operating AI platforms at enterprise scale, including monitoring for model drift, retrieval degradation, safety issues, latency, failed tool calls, and user feedback loops.

We track statistical output distributions to mitigate semantic drift. For RAG systems, Team C5MI monitors vector database retrieval health via context relevance metrics. Every transaction undergoes synchronous safety classification to intercept toxic inputs or jailbreak attempts. We track agent tool-calling success rates; spikes in timeouts trigger automated alerts and backup exception handlers. Finally, we integrate implicit/explicit feedback loops into operational UIs to drive routine fine-tuning and prompt refinement.

35. Explain your company's approach to observability for AI and agentic systems, including logs, traces, evaluations, metrics, dashboards, and incident response.

Team C5MI implements distributed tracing recording the entire execution chain under a single transactional trace ID (query, memory state, reasoning trajectory, API payload, response). We centralize telemetry into dashboards for real-time visibility. We configure automated alert thresholds for technical and AI quality metrics. If an agent enters runaway loops, encounters permission failures, or crosses cost-quotas, the system triggers incident response workflows to halt the session and notify the DevSecOps team.

36. Describe your company's approach to minimizing downtime during rollout of AI platforms, model updates, prompt updates, agent updates, or architecture changes. Provide examples of production cutover durations and business impacts from similar deployments.

Team C5MI utilizes Blue-Green deployments and Canary releases to route production traffic only after health checks pass. Prompts and system instructions are decoupled from application code, allowing hot-swaps without restarting containers. Built on containerized architecture, standard asset transitions achieve a near-zero downtime profile. This preserves asset visibility and prevents data processing blockages, ensuring continuity for high-volume operations.

37. Explain how your company manages ongoing platform sustainment, model lifecycle management, prompt maintenance, security patching, dependency updates, and technical debt reduction.

We treat AI infrastructure as a living ecosystem. LLMOps pipelines monitor prompts for semantic drift and accuracy degradation. Prompts are treated as code, tracked via version control, and optimized using user feedback. Team C5MI deploys automated vulnerability scanners to vet OS patches and software dependencies in staging before promotion. Dedicated engineering cycles address technical debt, including refactoring redundant agent paths, deprecating legacy model endpoints, and optimizing vector database indexes.

3.7 Implementation and Risk

38. Describe the implementation methodology or methodologies your company would recommend for executing these requirements, such as Agile, DevSecOps, phased rollout, pilot-to-production, product operating model, or hybrid approaches.

Team C5MI recommends a Hybrid Product Operating Model blending Scaled Agile with DevSecOps guardrails, delivered through a Pilot-to-Production lifecycle. We organize delivery around functional product teams working in multi-week sprints. Security and quality are "shifted left," requiring automated validation before moving out of development. We initiate deployments via MVP/Pilot groups, integrating lessons learned into the core baseline before scaling incrementally across the enterprise.

39. Describe your company's experience using your preferred methodology in large Federal AI, cloud modernization, or enterprise transformation efforts.

Team C5MI utilizes continuous CI/CD pipelines to automate the testing, validation, and deployment of complex data structures. In coordination with Google Public Sector, we have successfully integrated these practices for enterprise-wide transformation. By maintaining active stakeholder engagement and structured sprint reviews, we align capabilities with mission requirements while minimizing technical debt and accelerating the delivery of secure, accredited software solutions.

40. Given your understanding of a Government environment with multiple stakeholders, legacy systems, sensitive data, and evolving requirements, what is your typical implementation model for standing up Cloud-based LLM and agentic AI capabilities?

We recommend a parallel effort between data architecture, security, and user adoption teams. Team C5MI identifies data formats and owners while simultaneously engaging Government security to map blockers. User adoption teams conduct recurring touchpoints to incorporate feedback interactively. This collaborative structure ensures technical development remains aligned with user needs and security mandates, securing high user adoption from day one.

41. What is your estimated timeline from project initiation to pilot, initial operational capability, and full deployment? Please identify key milestones, dependencies, decision gates, and estimated durations.

Team C5MI estimates project initiation to pilot (MVP) takes 30–90 days. Initial Operational Capability (IOC) is achievable in 6–12 months, followed by full deployment within 12–24 months, depending on the scale and complexity of the integration. Key milestones include initial data ingestion, security ATO, and pilot user feedback integration. Decision gates are tied to successful performance metrics (e.g., accuracy, adoption) at the end of each pilot phase.

Team C5MI recommends a phased implementation approach that delivers value quickly while reducing risk. Project initiation, planning, and data onboarding are expected within the first 30–90 days, with initial operational capability achievable within 6–12 months depending on data availability, security approvals, and integration requirements. The effort begins with project planning, data access, and integration with DLA data sources, followed by the establishment of the cloud infrastructure, security controls, and DevSecOps pipelines. An initial MVP is then developed and connected to DLA data sources for testing and user feedback. Following user acceptance testing and refinement, the solution is deployed incrementally to user communities with supporting training, adoption activities, and AI champion development. Full deployment is expected within 12–24 months, including enterprise-wide rollout, user enablement, ongoing enhancements, and long-term sustainment. Key decision points throughout the effort include approval of project plans, security architecture, data access, MVP readiness, user testing results, and phased deployment milestones.

42. Describe your recommended sequencing for use-case selection, platform foundation, data onboarding, security controls, evaluation baselines, pilot execution, user adoption, and scaled deployment.

We recommend a strict, overlapping sequencing pattern. Data onboarding begins immediately, as it is a foundational dependency. Simultaneously, we conduct grassroots research to identify high-impact use cases. We establish the platform foundation after data evaluation, building security controls into the infrastructure from the outset. We then define evaluation baselines (task completion, accuracy) before executing a pilot period to collect feedback, finally rolling out scaled deployment team-by-team. Team C5MI recommends a concurrent, user-centered implementation approach that begins with immediate data onboarding while identifying high-impact AI use cases across the DLA enterprise. Security controls are incorporated from the start, and the platform foundation is established once data requirements and architecture are fully understood. Working closely with DLA stakeholders, the team defines success metrics such as efficiency gains, cost savings, accuracy, and process improvements to measure outcomes. An MVP is then piloted with select users to validate functionality, gather feedback, and refine capabilities before broader deployment. User adoption is embedded throughout the lifecycle through workforce training, AI Champions, bootcamps, brainstorming workshops, office hours, adoption analytics, and continuous feedback mechanisms to drive sustainable enterprise-wide adoption and mission impact. Following pilot testing and user feedback, the solution is deployed incrementally across DLA organizations to ensure successful adoption and sustained usage. This phased approach prioritizes data readiness, security, and measurable outcomes before expanding capabilities enterprise-wide.

43. What are the three to five major risks associated with your proposed approach, and what specific mitigation strategies would you apply to each?

1. Data Bottlenecks:
 - Risks: Timely access to required DLA data may impact AI solution development and deployment schedules.
 - Mitigation: Team C5MI can leverage synthetic data and proven data modeling approaches, drawing on extensive experience supporting DLA and other Federal agencies, to accelerate development while operational data access is being established.
2. Security Delays:
 - Risks: Required security reviews and ATO activities could impact implementation timelines.
 - Mitigated by leveraging existing GAILFORCE ATO architecture and Google Cloud's pre-accredited security frameworks.
3. Incremental Funding:
 - Risk: Projects funded in phases may face schedule uncertainty if priorities or funding levels change over time.
 - Mitigation: Mitigated by tying the AI project's funding to specific phases to ensure incremental gains are measurable and maintainable even if funding trajectories fluctuate.

44. Identify the Government actions, decisions, data access, infrastructure prerequisites, and policy inputs that would be required to enable successful implementation.

The Government must identify necessary data and engage Data Owners prior to contract kickoff. They must explicitly outline security requirements (ATO, ICAM role mappings, prompt-policy guidelines) from the outset. Finally, the Government must define the funding path to tie appropriate resources and milestones to each stage, preventing delays and ensuring a smooth, secure deployment.

3.8 Workforce, Pricing, and Delivery Model

45. What major labor categories would your company propose to execute these requirements? Include the role each labor category would perform in architecture, data engineering, security, responsible AI, model engineering, agent engineering, testing, integration, and sustainment.

Team C5MI proposes a scalable resource model scaled to mission requirements. Platform setup is resource-intensive in the first 90 days, focusing on DevSecOps and Data Engineering. Pilot implementation shifts resources toward Model/Agent Engineering and QA. Production hardening and sustainment emphasize maintenance, monitoring, and continuous integration, requiring a steady-state team of AI practitioners and security specialists.

- **Technical Project Lead:** Provides overall technical leadership and architectural oversight, ensuring the AI platform is scalable, secure, and interoperable across hybrid and multi-cloud environments.
- **Enterprise AI/Cloud Architect:** Leads overarching hybrid/multi-cloud systems architecture design, ensuring cross-platform interoperability and zero vendor lock-in.
- **Data Engineer:** Manages structured and unstructured ingest pipelines, structural document chunking, automated metadata tagging, and end-to-end data lineage tracking.
- **DevSecOps Security Engineer:** Implements perimeter boundaries, logical workload isolation, key management (CMEK), network isolation, and continuous ATO evidence compilation.
- **Model & Agent Engineer:** Conducts model evaluation, parameter-efficient fine-tuning (PEFT), tool schema configuration, and designs collaborative multi-agent architectures.
- **AI Software Engineers:** Develop, integrate, and optimize AI/ML capabilities, including model evaluation, fine-tuning, agent orchestration, and application development to meet mission requirements.
- **UI/UX Engineer:** Designs intuitive, accessible, and mission-focused user experiences that align with DLA branding, usability standards, and operational workflows.
- **QA / Independent Testing Specialist:** Conducts adversarial red-teaming, automated regression testing, and metrics tracking (grounding fidelity, latency, and TTFT).
- **Responsible AI & Safety Specialist:** Sets up input/output guardrail classifiers, safety content moderation filters, systemic bias testing, and standard escalation procedures.
- **Sustainment & MLOps Platform Specialist:** Oversees automated monitoring for model drift, retrieval degradation, runtime exception handling, and platform cost-quota optimization.
- **AI Adoption & Engagement Lead:** Leads training, user engagement, and change management activities to ensure successful adoption, effective utilization, and long-term value realization of the AI platform.

46. Provide estimated labor hours or level-of-effort ranges by labor category for platform setup, pilot implementation, production hardening, rollout, and sustainment.

To meet the Statement of Work (SOW) requirements outlined in the RFI, Team C5MI proposes a dedicated team of 12 personnel. The table below details the labor categories and annual Level of Effort (LOE) required to support platform development, security enhancement, long-term sustainment, and successful user adoption. /Year

Labor Category	Estimated Hours
Technical Project Lead	1920
Enterprise AI/Cloud Architect	1920
Data Engineer	7000

DevSecOps Security Engineer	1920
Model & Agent Engineer	4160
AI Software Engineers	5760
UI/UX Engineer	1920
QA / Independent Testing Specialist	1920
Sustainment & MLOps Platform Specialist	1920
AI Adoption & Engagement Lead	1920

47. Describe your company's recommended delivery model, including on-site versus remote support, product team composition, surge support, training, and knowledge transfer to Government personnel.

Team C5MI recommends an agile, outcome-focused Product Team Delivery Model. We combine targeted on-site engagement for architecture mapping with secure remote engineering for core development. To handle surge requirements, we utilize an elastic bench of pre-vetted specialists. We embed knowledge transfer into every sprint through pair-programming and workshops, ensuring Government personnel are prepared for long-term operations.

48. Describe your company's assumptions regarding on-site and remote work for this effort. Please state whether your estimated level of effort and pricing model includes travel to DLA locations (e.g., Fort Belvoir, VA) and, if so, the anticipated frequency and duration of such travel.

Team C5MI assumes 80% remote and 20% on-site work. Our pricing model explicitly includes travel costs to primary DLA locations, including Fort Belvoir, VA. We anticipate a cadence of one week per month (3–5 days) for key personnel, strategically timed for sprint planning cycles and technical program reviews.

49. Describe your company's approach to enabling Government self-sufficiency over time, including training, documentation, prompt governance, admin controls, platform operations, and transition planning.

We deliver living technical documentation, including IaC scripts and API catalogs. Team C5MI deploys centralized administrative dashboards for Government personnel to control prompt libraries and RBAC schemas. We implement a graduated transition methodology, shifting from a contractor-led model to a co-engineering paradigm where DLA personnel drive platform operations, ensuring Government self-sufficiency.

3.9 Optional Government-Specific Questions

50. Describe your company's experience building AI assistants or agents that support logistics, supply chain, cataloging, maintenance, acquisition, customer service, or enterprise operations.

Team C5MI has a proven track record building assistants that automate asset tracking and supply chain queries. We have deployed conversational assistants allowing maintenance personnel to query material catalogs and verify stock availability. Our agents interpret error tickets, query historical troubleshooting workflows, and draft remediations. We also engineer automated classification models that parse procurement files, isolate transactional discrepancies, and map logs to enterprise databases, significantly reducing backlogs.

51. Describe your experience using LLMs or agents to work with highly structured authoritative data together with large volumes of unstructured documents, manuals, policy artifacts, and historical records.

We bridge the gap between transactional databases (SAP) and fragmented document repositories. Our ingestion pipelines process structured enterprise info alongside unstructured files (manuals, policy). Using vector embedding models and metadata mapping, we ground LLMs across these datasets. Agents extract transactional status history while retrieving troubleshooting guidelines from unstructured manuals. RAG architectures force the model to provide explicit inline citations and direct hyperlinks back to both database rows and document sections, ensuring outputs are fully traceable.

52. Describe your experience implementing human-reviewed AI workflows where outputs inform, but do not automatically execute, mission or business decisions.

Team C5MI designs applications in strict adherence to RAI guidelines, enforcing that AI serves as a decision-support capability. We hardcode execution checkpoints into workflows, programmatically blocking autonomous finalization. The system packages findings, reasoning trajectories, and citations into a staging UI, requiring an authorized operator to explicitly confirm, modify, or reject the output before execution. This ensures human accountability remains the core driver.

53. Describe your experience implementing agentic AI solutions that interact with ERP, workflow, or ticketing systems while preserving approval chains, transaction integrity, and auditability.

We integrate agentic solutions with platforms like SAP and ServiceNow while preserving integrity. Agents interact via constrained API gateways (e.g., Apigee) utilizing OAuth scopes. Team C5MI implements secure gateways (OData/BAPI) to guarantee atomic database transactions. Every action is logged with transactional IDs in an immutable ledger. For operations carrying financial weight, the agent generates an approval checkpoint, ensuring no automation bypasses established business rules.

54. How does your firm define Agentic AI vs. standard LLMs, and what security boundaries do you recommend the Government put in place?

Team C5MI defines Agentic AI as a goal-oriented, stateful orchestration wrapper equipped with reasoning frameworks, memory management, and access to external execution tools distinct from passive, state-free LLMs. We recommend sandboxing all tool execution through a monitored API management tier using rigid schemas. We also recommend administrative iteration caps to eliminate runaway execution loops.

55. Review the attached Draft LCATs. Do you believe these education and experience requirements reflect current commercial AI industry standards? If not, what do you recommend?

Yes, we do. While baseline requirements are strong, we recommend targeted updates. Team C5MI suggests adding explicit experience with Parameter-Efficient Fine-Tuning (PEFT), LoRA, and structured schema validation for Model/Agent Engineers. The Responsible AI Specialist should explicitly require background in bias testing and injection mitigation. Legacy system descriptions should be updated to focus on modern LLMops, including monitoring statistical model drift.

56. Given potential DoD restrictions on certain commercial models (e.g., Anthropic), how would your architecture remain model-agnostic and ensure continuity of operations?

Team C5MI treats foundation models as interchangeable engines via unified abstraction layers. We decouple reasoning logic, prompt registries, and RAG pipelines from provider-specific APIs, allowing for hot-swaps via configuration rather than code rewrites. Our architecture routes exclusively to authorized model registries. If a restriction occurs, the system supports rapid migration to enterprise-grade open models running locally within containerized environments, ensuring DLA operations remain operational. This capability allows us to fully

abide by DoD restrictions around (e.g. Anthropic) and only select models that are approved and certified for DoD and the DLA.

4.0 Recommendations

Based on our experience delivering enterprise AI and analytics capabilities in Federal environments, we offer the following recommendations for DLA's consideration as the acquisition strategy is refined.

Performance-Based Metrics: We recommend objective, measurable standards for AI platform delivery and operations:

- **Model Quality:** Hallucination rate below 5% on grounding fidelity tests; citation accuracy above 95% for all RAG responses
- **Platform Reliability:** 99.5% uptime for LLM endpoints and vector databases; mean time to resolution under 4 hours for critical incidents
- **Response Latency:** Sub-3-second response time for standard RAG queries; sub-10-second completion for multi-step agentic workflows
- **Security Compliance:** Zero unauthorized data exfiltration events; 100% audit log completeness for all model calls, prompts, and agent actions
- **Knowledge Transfer:** Government team independently executes prompt updates, model evaluations, and deployment rollbacks within 90 days of transition

Performance-based standards ensure DLA pays for measurable outcomes rather than labor hours. This approach enables objective CPAR evaluation and incentivizes continuous improvement throughout the contract period.

Technology Validation Requirements

DLA should require offerors to demonstrate AI capabilities in a controlled environment as part of the solicitation evaluation process. Validation demonstrations should include live RAG pipeline queries against sample DLA data, agentic AI workflow execution with human-in-the-loop approvals and rollback scenarios, and prompt injection testing to verify guardrail effectiveness. This approach reduces risk by ensuring the successful offeror can deliver functional capabilities before enterprise deployment. There is no AI without data; participants should demonstrate ability to navigate complex DLA datasets as they are structured out of DLAs critical ERP systems.

Integrated Service Delivery Model

Given the interdependencies across AI model engineering, data pipelines, security controls, and enterprise integration, DLA should structure task orders to enable cross-functional product teams rather than siloed technical disciplines. Our delivery model assigns dedicated teams that span architecture, data engineering, MLOps, security, and integration roles. Each team owns end-to-end delivery for specific use cases or capability areas. This structure eliminates coordination delays between separate contractors and creates shared accountability for operational outcomes. Integrated teams respond faster to evolving requirements and reduce handoff risks during production deployment.

Contract Vehicle

C5MI is a current JETS 2.0 IDIQ holder (Group A, Small Business) under CAGE Code 9ALF5. JETS 2.0 supports IT services, technical expertise, and lifecycle management across DLA's IT Enterprise, including application development, data analytics, and administrative program management support.

The JETS 2.0 vehicle provides DLA with proven mechanisms for small business competition (Group A IDIQ holders), 8(a) competition (Group B awardees), and full competition (Group C for other-than-small business). Given that AI platform services align with JETS 2.0 scope under NAICS 541611 (Administrative Management and General Management Consulting Services) and NAICS 541512 (Computer Systems Design Services), leveraging JETS 2.0

as a single award, Small Business set-aside will reduce acquisition timeline and support DLA's socioeconomic goals.

Alternative vehicles include GSA MAS (less tailored to DLA-specific AI requirements) or a standalone IDIQ (longer acquisition timeline, higher administrative burden).

Contract Type

Firm-Fixed-Price (FFP) task orders incentivize efficiency and provide cost predictability. For AI services with defined outcomes (pilot delivery within 60 days, 95% citation accuracy, 99.5% platform uptime), FFP ensures DLA pays for results rather than labor hours. Performance risk transfers to the contractor, which aligns incentives with mission success.

For early exploratory work or evolving requirements, Time and Materials (T&M) task orders provide flexibility during discovery and prototyping phases. A hybrid approach using T&M for initial use case definition and FFP for production deployment balances flexibility with cost control.

Source Selection Method

Oral presentations combined with written technical proposals enable DLA to assess vendor understanding beyond paper submissions. For AI capabilities, oral presentations allow evaluators to probe RAG architecture design decisions, model evaluation methodologies, and agentic AI guardrail implementation in real time.

Live technical demonstrations should accompany oral presentations. Demonstrations validate that vendors possess functional AI capabilities rather than theoretical approaches. For example, vendors demonstrate a working RAG pipeline retrieving data from a sample document set, an agent executing a multi-step workflow with rollback capability, and automated evaluation tests measuring hallucination rates.

This approach reduces technical evaluation risk and ensures selected vendors can execute on day one.

Labor Category Flexibility

The rapid evolution of AI technologies requires labor category structures that accommodate emerging roles. DLA should allow contractors to propose labor category modifications during the contract period to address new AI disciplines (LLMOps specialists, Responsible AI engineers, adversarial testing experts) as commercial standards evolve.

Ceiling rates for new labor categories should be negotiated based on fair and reasonable price analysis per FAR 15.404-1(b)(2), with modifications executed before personnel assignment. This approach ensures DLA accesses current market talent without requiring full contract recompetes.

Data Access and Infrastructure Readiness

Successful AI implementation depends on timely Government decisions regarding data access, infrastructure provisioning, and policy inputs. DLA should plan for dedicated Government Product Owner availability, data access approvals within 30 days of task order award, accredited cloud infrastructure (Impact Level 4 or higher) available at contract start, and documented prompt policy guidelines defining prohibited vocabulary and operational boundaries.

Delays in these enablers directly impact delivery timelines. Early coordination between DLA technical teams, security offices, and contracting personnel reduces schedule risk.

These recommendations reflect lessons learned from our current DLA platform delivery experience and aim to strengthen the acquisition strategy and delivery of AI capabilities to DLA. C5MI welcomes the opportunity to discuss these concepts with DLA leadership as acquisition planning matures.

5.0 Conclusion

We thank you again for including Team C5MI in this RFI and for reviewing our response. Team C5MI's extensive experience in data analysis, data visualization, and project management, combined with our proven track record of successful software implementations and deep knowledge of the DLA's current technological landscape, position us as a strong partner to support the DLA in developing congressional and legislative action reports. We are committed to providing accurate, timely, and data-centric reporting to inform DLA leadership and operations as it pertains to legislative actions. We welcome any additional questions you may have, and we look forward to the forthcoming Request for Proposal (RFP).